

Platform Autoscaling Dev Demo

It's here.

Troy Sankey

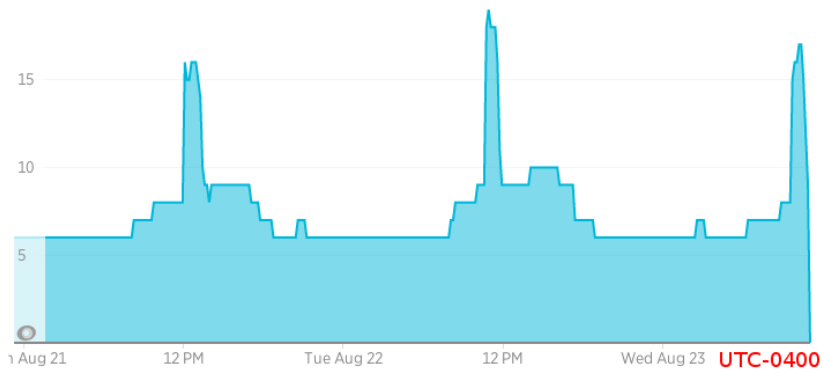
August 25, 2017

Configuration

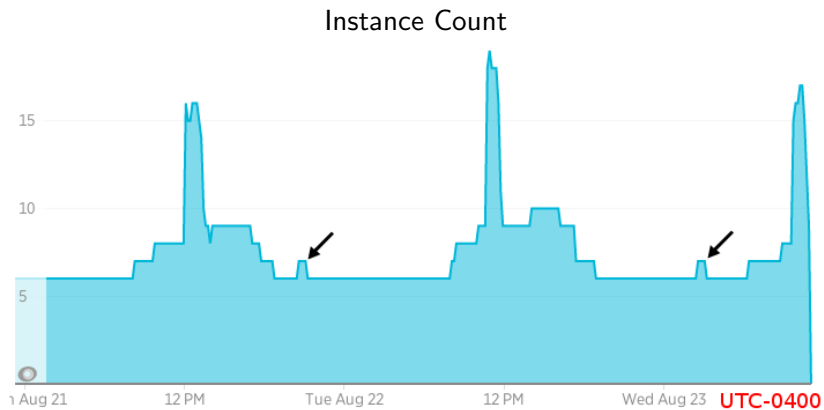
- ▶ Initially, the instance count was 12, now we scale between 6 and 12.
- ▶ Average CPUUtilization is controlled to stay between 35% and 55%.
- ▶ 10 minute instance launch delay.

Stability

Instance Count

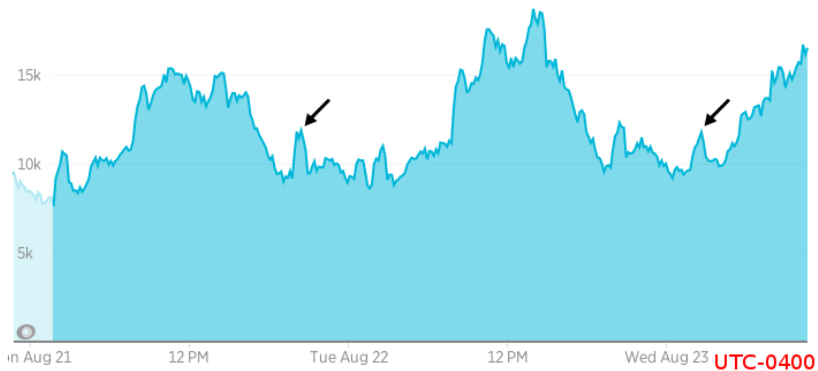


Stability



Stability

Throughput (requests/minute)



Cost

- ▶ Average of 6.8 instance over the last week
- ▶ Initial yearly cost savings of about \$10,000, after adjusting reservations from 12 to 6 instances.

- ▶ About 4ms increase in median response time (+10%).